

Capital Chronicle ContentOps — Pre-Alpha General/Process Content Master Plan v1

0. Operating Decision

Capital Chronicle ContentOps should not pause completely during the internal-alpha wait-state.

The correct strategy is:

Maintain the artifact-backed wait-state, but open a tiny manual pre-alpha general/process content lane.

This lane exists to build audience, trust, and product positioning before Capital Chronicle has approved internal-alpha artifacts. It must not produce fake market calls, fake artifact-backed content, synthetic performance content, public-postable fixture content, or anything that implies Capital Chronicle already has live forecast authority.

The repo-side work should begin with policy, taxonomy, QA gates, and negative tests — not a content generator and not a publishing/export workflow.

1. Current Authority Baseline

Repo: `A:\Capital Chronicle\tools\cc-live-contentops`

Current accepted state:

- ContentOps is in terminal local alpha wait-state.
- Real Capital Chronicle alpha artifacts do not exist in this sidecar yet.
- The existing stack is local-only, deterministic, fixture-only, and human-review-required.
- Provider/LLM API calls, network/search, platform APIs, credentials/env reads, scheduling, live posting, autonomous replies/DMs, browser automation/scraping, public-postable synthetic content, real alpha artifact access, and core repo reads/writes remain disabled.
- The next pointer remains:

`WAIT_FOR_REAL_CAPITAL_CHRONICLE_ALPHA_ARTIFACTS_OR_OPERATOR_SELECTED_LOCAL_MAINTENANCE`

2. Master Rule

Content maturity cannot outrun data maturity.

Until real approved Capital Chronicle alpha artifacts exist, ContentOps may support only:

- general/process policy;
- taxonomy;

- editorial QA;
- negative guardrail tests;
- operator checklists;
- non-public dry-run previews;
- documentation and evidence updates.

It must not generate public-ready fake posts from fixtures or imply artifact-backed output.

3. Two-Lane Content Model

Lane A — Pre-Alpha General/Process Content

Status: allowed conceptually, but repo support must start with policy/QA only.

Definition: General/process content is first-party educational or philosophical content about:

- why Capital Chronicle is being built;
- how serious macro research should handle missing data;
- why forecast readiness matters;
- why no-forecast is often the correct answer;
- how bad forecasts fail;
- why the product is not a signal service.

This lane does not require Capital Chronicle alpha artifacts.

Allowed evidence basis:

- Jim's own product philosophy;
- evergreen macro concepts;
- public official sources if a factual claim needs citation;
- clearly labeled general frameworks;
- process notes that do not imply live system output.

Never allowed:

- source artifact IDs invented for credibility;
- DQR/forecast-readiness states invented as if produced by Capital Chronicle;
- fake lineage;
- fixture/demo content presented as real;
- market calls;
- predictions framed as actionable;
- performance claims.

Lane B — Future Artifact-Backed Capital Chronicle Content

Status: blocked until real approved alpha artifacts exist.

Definition: Artifact-backed content depends on real Capital Chronicle outputs and must include:

- operator-approved artifact spec;
- source artifact IDs;
- lineage references;
- freshness;
- limitations;
- DQR/data sufficiency/forecast readiness status;
- missing/proxy/degraded labels;
- content type mapping;
- no financial advice or signal/execution language.

This lane resumes only through the existing real-artifact intake gate and must route to `READY_FOR_LOCAL_REVIEW_ONLY`, never directly to public-ready.

4. Safe Pre-Alpha Content Pillars

4.1 Build-in-Public: Process, Not Results

Purpose: Build trust by showing the discipline of refusing weak evidence.

Safe themes:

- why the system should block forecasts when data is insufficient;
- why manual review exists;
- why local-first and human-in-the-loop matter;
- what the team is building before alpha;
- what the system is designed not to do.

Example angles:

- “Why Capital Chronicle would rather publish no forecast than a low-quality forecast.”
- “The hardest part of macro research is not prediction; it is knowing when evidence is not enough.”
- “A good research system needs a refusal mode.”

4.2 Macro Education: Evergreen, Non-Directional

Purpose: Teach macro reasoning without telling anyone what to trade.

Safe themes:

- data revisions;
- official-source lag;
- soft vs hard data;
- base effects;
- cross-asset confirmation;
- why one print rarely settles a macro thesis;

- difference between observation, thesis, and actionable conclusion.

Example angles:

- “A single CPI print is not a regime model.”
- “Why macro data can be true, stale, and still misleading.”
- “The difference between a chart and a forecast.”

4.3 Data Sufficiency / Forecast Readiness Education

Purpose: Make Capital Chronicle's wedge legible before alpha.

Safe themes:

- missing stays missing;
- degraded stays degraded;
- proxy-only is not official truth;
- confidence is invalid if required inputs are absent;
- forecast readiness is a gate, not a marketing badge.

Example angles:

- “What has to be true before a forecast is even worth drafting?”
- “Why proxy-only data should be labeled, not hidden.”
- “No forecast is sometimes the most honest output.”

4.4 Failure Forensics Philosophy

Purpose: Prepare the audience for the later forensics product without needing real outcomes yet.

Safe themes:

- stale-data error;
- missing-lane error;
- regime-classification error;
- timing error;
- catalyst error;
- overconfidence error;
- source-quality error.

Example angles:

- “Bad forecasts are not always bad models. Sometimes the system did not know what it was missing.”
- “A forecast can fail because the thesis was wrong, the timing was wrong, or the input fabric was incomplete.”
- “Why post-move forensics may be more useful than pre-move confidence.”

Pre-alpha limitation: Avoid detailed public post-mortems of live market events unless they are clearly evergreen education and source-cited. Do not turn them into implied trade lessons.

4.5 Product Philosophy / Positioning

Purpose: Define what Capital Chronicle is and is not.

Safe themes:

- not a Bloomberg replacement;
- not an AI trading bot;
- not a signal service;
- not an execution engine;
- not guaranteed forecasting;
- a local-first macro research system focused on data sufficiency, forecast readiness, and failure forensics.

Example angles:

- “Capital Chronicle is not built to shout signals. It is built to ask whether a thesis deserves confidence.”
- “The wedge is not speed. The wedge is evidence discipline.”
- “The product is designed to make uncertainty visible.”

5. Forbidden Content Categories

Always block:

- buy/sell/hold language;
- position sizing;
- entry/exit levels;
- personalized financial advice;
- “watch this level” as trade implication;
- “my model says buy/sell/short/long”;
- price targets;
- guaranteed or near-guaranteed prediction;
- probability claims about future market moves unless backed by a future approved artifact and reviewed policy;
- “this is not financial advice, but...” followed by actionable framing;
- broker/order/execution/MT5/live trading language;
- signal-service positioning;
- fake performance;
- simulated or fixture performance presented as public proof;
- synthetic/demo fixture content made public-postable;
- raw vendor data redistribution;
- unverified numeric market claims;
- public claim that Capital Chronicle alpha is already active before approved artifacts exist;
- content that hides missing/degraded/proxy data;
- content that converts blocked readiness into confidence.

6. Platform Strategy

Substack

Role: Canonical long-form home.

Use for:

- macro education essays;
- product philosophy letters;
- failure-forensics frameworks;
- data sufficiency explainers;
- public “why we wait” essays.

Pre-alpha cadence: One essay every one to two weeks.

Rules:

- no live market calls;
- no performance claims;
- no fake artifact-backed sections;
- every factual/current claim must have source notes;
- general/process content must be labeled as general/product/process, not Capital Chronicle output.

LinkedIn

Role: Professional positioning and founder/operator voice.

Use for:

- concise essays;
- build-in-public notes;
- process discipline;
- product philosophy;
- “why this system refuses weak forecasts.”

Pre-alpha cadence: One to two posts per week.

Rules:

- professional, calm, explicit limitations;
- no hype;
- no “AI trading bot” framing;
- no directional market implication;
- avoid current-market commentary unless source-cited and educational.

X

Role: Sharp hook and conversation layer.

Use for:

- short principles;
- threads derived from Substack essays;
- build-in-public snapshots;
- “no forecast” philosophy.

Pre-alpha cadence: Optional and conservative.

Rules:

- avoid hot takes;
- avoid live-event market interpretation;
- never thread a quasi-signal;
- every post should be able to stand as education/process, not advice.

Threads

Role: Lower-stakes community mirror.

Use for:

- softer educational fragments;
- build-in-public notes;
- simple questions around process;
- less technical versions of LinkedIn/X ideas.

Pre-alpha cadence: Optional.

Rules:

- no market calls;
- no reactive political/market takes;
- use as community tone testing, not authority channel.

7. Manual-Only Editorial Workflow

Pre-alpha content should move through five gates:

Gate 1 — Lane Classification

Every draft must be one of:

- `pre_alpha_general_process`
- `future_artifact_backed`

During wait-state, only `pre_alpha_general_process` may be discussed for public manual posting outside the repo. The repo should not produce public-ready content until the QA lane exists.

Gate 2 — Claim Classification

Every draft claim must be classified as:

- first-party philosophy;
- evergreen education;
- cited factual claim;
- current factual claim requiring citation;
- market-sensitive claim;
- forbidden claim.

Market-sensitive claims are blocked unless transformed into evergreen education.

Gate 3 — Source Requirement

If the claim is factual and not common/evergreen, it needs:

- source;
- date;
- limitation;
- freshness note if relevant.

If a source is missing, the draft is not publishable.

Gate 4 — Forbidden Language Scan

Block:

- trade/action verbs;
- certainty/guarantee terms;
- signal-service wording;
- broker/order/execution wording;
- fake performance;
- implied alpha output;
- unsupported market direction.

Gate 5 — Human Final Review

Only Jim can decide to publish manually.

ContentOps must not:

- approve;
- auto-select final public copy;
- post;
- schedule;
- send platform API requests;
- scrape metrics;
- read credentials;
- DM/reply autonomously.

8. Recommended Next Repo Backlog

TASK_CONTENTOPS_0075_PRE_ALPHA_GENERAL_PROCESS_POLICY_AND_QA_GATE_V0

Objective: Create a local-only policy/taxonomy/QA gate for pre-alpha general/process content without generating public-ready content.

Allowed:

- docs;
- schema/frontmatter contract;
- deterministic QA checklist;
- negative tests;
- forbidden language test fixtures;
- read-only/dry-run CLI summary if useful;
- evidence packet.

Forbidden:

- no content generator;
- no public-postable outputs;
- no export/clipboard bundle;
- no platform templates that look publish-ready;
- no provider/search/platform/API/credential access;
- no real alpha artifact access;
- no core repo reads/writes;
- no scheduling;
- no auto-approval.

Acceptance:

- `pre_alpha_general_process` is clearly separated from future `artifact_backed`.

- General/process content cannot claim artifact lineage.
- General/process content cannot route to the real artifact intake/bridge path.
- Market call examples are blocked.
- Signal-service/execution/broker examples are blocked.
- Fixture/demo content remains not-public-postable.
- Terminal wait-state pointer is preserved.
- Full tests pass.
- Suspicious scan passes.
- Evidence packet confirms no network/provider/search/platform/API/credential/live action.

TASK_CONTENTOPS_0076_PRE_ALPHA_EDITORIAL_REVIEW_PACKET_DRY_RUN_V0

Only after 0075 PASS.

Objective: Create a non-public, local-only editorial review packet format for Jim-authored general/process drafts.

Allowed:

- manual input fixture representing Jim-authored text;
- QA result;
- claim classification;
- source-needed flags;
- platform-fit notes;
- blocked-language report;
- “not public-postable until Jim final review” status.

Forbidden:

- no synthetic public-ready content;
- no automated final copy;
- no platform export;
- no API;
- no credentials;
- no posting.

Acceptance:

- review packet helps Jim decide what is safe to manually rewrite/post;
- output is explicitly review-only;
- no draft is marked publish-ready by the system.

TASK_CONTENTOPS_0077_PRE_ALPHA_MANUAL_AUTHORSHIP_OPERATING_GUIDE_V0

Only after 0076 PASS.

Objective: Add a concise operator guide for how Jim can write and manually publish pre-alpha general/process content outside the repo while keeping ContentOps records clean.

Allowed:

- docs-only guide;
- examples as non-public sample ideas, not ready-to-post copy;
- checklist;
- manual log template.

Forbidden:

- no generated social posts;
- no automated export;
- no platform integrations;
- no metrics ingestion;
- no scheduling.

Acceptance:

- Jim has a clear manual workflow.
- The repo remains local-only and not public-postable.
- The guide explains the difference between manually authored public content and repo-generated fixture/dry-run content.

TASK_CONTENTOPS_0078_GROUNDED_RESEARCH_BRIEF_TEMPLATE_DESIGN_ONLY_V0

Optional, after Deep Research findings are finalized.

Objective: Create a design-only template for future grounded research briefs that can inform editorial strategy without becoming source authority or platform automation.

Allowed:

- template docs;
- source taxonomy;
- claim-risk taxonomy;
- "research context only" labels.

Forbidden:

- no search/provider integration in repo;
- no fetched sources;
- no live calls;
- no external credentials.

Acceptance:

- future research can be attached manually as context;
- deterministic QA still decides whether claims are allowed.

WAIT_FOR_REAL_ALPHA_ARTIFACTS_RESUME_PATH

When real Capital Chronicle internal-alpha artifacts exist, do not use the pre-alpha lane to ingest them.

Resume via the existing artifact-backed path: 1. Confirm artifact spec. 2. Confirm approved export location. 3. Confirm source artifact IDs. 4. Confirm lineage, freshness, limitations. 5. Confirm DQR/data sufficiency/forecast readiness states. 6. Confirm missing/proxy/degraded labels. 7. Confirm content type. 8. Confirm no financial advice/execution/signal claims. 9. Intake gate routes to `READY_FOR_LOCAL_REVIEW_ONLY`. 10. Jim reviews final copy manually before any public use.

9. What Not To Build Yet

Do not build:

- live publisher;
- platform adapters;
- scheduler;
- auto-cross-poster;
- API clients;
- credential storage;
- browser automation;
- scraping;
- autonomous replies/DMs;
- analytics fetcher;
- public-ready synthetic content generator;
- fake performance dashboard;
- market commentary generator;
- current-event take generator;
- “alpha teaser” generator.

These are premature and violate the wait-state risk posture.

10. Decision Tree

If Jim wants to start audience-building now

Proceed with 0075 first: Policy + taxonomy + QA gate only.

Jim may separately write public posts manually using the safe pillars, but repo outputs remain non-public unless and until later gates are built.

If Jim wants maximum safety and zero public presence

Do no content-lane implementation. Only maintain repo docs/tests and wait for real alpha artifacts.

If real alpha artifacts arrive soon

Do not expand pre-alpha lane. Switch to artifact-backed resume checklist and run the first real artifact intake task.

If Deep Research produces stronger platform guidance

Fold it into docs as advisory strategy only. Do not implement platform-specific automation.

11. Master Recommendation

Proceed with a small pre-alpha content lane, but implement it in this order:

1. **0075 — policy/taxonomy/QA barrier**
2. **0076 — review-only packet for Jim-authored drafts**
3. **0077 — manual authorship operating guide**
4. **0078 — optional research brief template**
5. **Real alpha resume path when artifacts exist**

The key principle is:

Build the guardrails before the generator. Build the lane boundary before the content workflow. Keep all public posting manual. Keep artifact-backed content blocked until real alpha artifacts exist.