

Addendum — Grounded News / Research Context Lane

Purpose

During the pre-alpha wait-state, Capital Chronicle ContentOps should support a safe real-content workflow using LLM-assisted grounded research and current news context.

The goal is to create timely, engaging social/newsletter content without pretending that Capital Chronicle has live alpha artifacts, trading signals, or forecast authority.

This lane uses real public information and current news as hooks, but turns them into education, process commentary, data-sufficiency analysis, forecast-readiness discussion, or product philosophy — not market calls.

Core Principle

News is a hook, not a signal.

A current event may be used to explain:

- why data quality matters;
- why a macro claim is not forecast-ready;
- how official sources should be checked;
- why markets often overreact to incomplete information;
- what uncertainty remains;
- what a serious research system would need before forming a thesis.

A current event must not be used to say:

- buy/sell/hold;
- long/short;
- target price;
- “this means X asset will move”;
- “our model predicts”;
- “this is the setup”;
- “watch this level”;
- “Capital Chronicle alpha says...”

Allowed Source Types

Preferred:

- official sources: central banks, statistical agencies, Treasury/fiscal agencies, regulators, exchanges when relevant;
- reputable news sources for event context;
- company/government primary releases when discussing policy or macro events;

- public research reports where redistribution is allowed;
- platform-native public posts only as context, not authority.

Avoid:

- unsourced social media rumors;
- anonymous market claims;
- raw vendor/proprietary data;
- paywalled content copied into outputs;
- screenshots without source/date;
- influencer claims as evidence;
- politically charged commentary without careful source handling.

Workflow

1. Operator or ChatGPT Deep Research collects a grounded research brief outside the repo.
2. The brief includes source URLs, publication dates, short source summaries, and claim-risk notes.
3. LLM proposes safe angles:
4. macro education angle;
5. data sufficiency angle;
6. forecast readiness angle;
7. failure forensics angle;
8. product philosophy angle;
9. build-in-public angle.
10. ContentOps QA checks:
11. forbidden trading/signal language;
12. unsupported numeric claims;
13. missing citations;
14. implied forecast authority;
15. artifact-backed claims without artifacts;
16. whether the post is educational/general/process content.
17. Jim manually approves, edits, and posts.
18. Any public URL/metrics are recorded manually only, if needed.

Content Types Added

Add or recognize these pre-alpha content subtypes:

- grounded_news_context
- official_data_explainer
- policy_process_commentary
- macro_education_from_news

- `forecast_readiness_from_news`
- `data_sufficiency_from_news`
- `failure_forensics_from_news`

These remain under the broader class:

`pre_alpha_general_process`

They are not artifact-backed Capital Chronicle outputs.

Safe Examples

Example 1: A central bank speech is released.

Unsafe: "The Fed is clearly preparing to cut. Duration longs look attractive."

Safe: "Central bank speeches are useful context, but they are not a forecast by themselves. A serious macro workflow still has to check inflation trend, labor-market cooling, financial conditions, and revision risk before treating the speech as thesis-ready."

Example 2: A CPI report is published.

Unsafe: "This CPI print means equities should rally."

Safe: "One CPI print can change the conversation, but it should not become a regime call by itself. The better question is whether the print confirms a broader disinflation path, whether revisions matter, and whether cross-asset confirmation exists."

Example 3: A geopolitical headline moves oil.

Unsafe: "Oil is going higher from here."

Safe: "Geopolitical headlines often create immediate price reaction before data sufficiency exists. A research process should separate headline risk, supply-chain evidence, inventory data, policy response, and positioning before calling it a durable macro shift."

Platform Use

Substack: Use for the full grounded explainer.

LinkedIn: Use for professional summary and process insight.

X: Use for concise hooks and short educational threads.

Threads: Use for softer, conversational versions.

All platforms: No signal framing. No urgency. No actionable trade language. No “model says” language. No fake alpha.

Repo-Side Implementation Sequence

Update the next backlog:

1. **0075 — Pre-Alpha Policy + QA Gate** Add `grounded_news_context` as a safe subtype under `pre_alpha_general_process`. Add negative tests proving news hooks cannot become market calls.
2. **0076 — Grounded Research Brief Schema** Local-only schema/template for operator-supplied research briefs. No network calls. No provider calls. No search integration. It only validates manually supplied source metadata and claim-risk labels.
3. **0077 — LLM-Assisted Draft Review Packet** Review-only packet for drafts produced outside the repo by Jim/ChatGPT. The system checks citations, risk language, source freshness labels, and content class. It does not auto-generate final public copy.
4. **0078 — Manual Publish + Metrics Capture Guide** Jim manually posts approved content. ContentOps may record URLs and manually observed metrics only. No scraping, no platform API, no credentials.

Final Rule

The grounded-news lane is allowed only if every item remains:

- source-cited;
- educational/general/process;
- manually reviewed;
- manually posted;
- non-advisory;
- non-signal;
- non-artifact-backed unless real approved artifacts exist.